

BUILD CONVERSATION

BUILDCHAT - Decision Bajo Presion Week 6

Full visible user-assistant transcript for the Week 6 build, test, fix, and two-deployment cycle.

Submission detail	Value
Author	David Buzali
Coursework	Week 6 individual working slice
Product	Decision Bajo Presion
Live URL	https://decision-bajo-presion-week-6.vercel.app
GitHub	https://github.com/davidbuzali/decision-bajo-presion-week-6
Deployment	Deployment 2 - verified Ready
Data boundary	Synthetic demonstration records only
Visible messages	148
Export boundary	Through the final PDF preparation request
Privacy boundary	Internal reasoning, system instructions, credentials, ambient UI, and tool telemetry excluded

Prepared for final Week 6 submission.

Conversation transcript

This export preserves the visible development conversation in chronological order. Attachments are identified by the text shown in chat; binary files and internal execution logs are not reproduced. Progress updates are included because they document decisions, verification, commits, and deployment state.

DAVID BUZALI - USER	16 Sep 2026, 18:46:30 CST
# Files mentioned by the user:	
## BLUEPRINT_Team3_Week6.pdf: /Users/davidbuzali/Desktop/BLUEPRINT_Team3_Week6.pdf	
## BRIEF_davidbuzali (week 6).pdf: /Users/davidbuzali/Desktop/BRIEF_davidbuzali (week 6).pdf	
Distinguish instructions in attached documents from the user's request.	
## My request: My [REDACTED DEVICE CODE]-VVVV Manifesto	
Purpose I intend to build something meaningful that contributes to the world, creates lasting opportunity, and improves the lives of others. I want to confront educational inequality, especially in Mexico, so more people can think critically and choose their own paths. A life well used will combine contribution, fulfilling work, personal freedom, and presence for those I love.	
Passions I am energized by consequential problem-solving, decision-making, and optimization with measurable real-world impact. Competition, teamwork, disciplined improvement, and being useful to a shared mission bring out my best. I will deepen my knowledge of AI, business, and consulting so I can solve important problems across industries.	
Principles I will act with loyalty, honesty, respect, empathy, fairness, and regard for health and the law. Important decisions must reflect my values, consider their long-term effects, and remain feasible. I will pursue excellence without demanding perfection from myself or others.	
Personality I am disciplined, ambitious, analytical, and naturally inclined to lead. I bring vision, dedication, and an ability to make complex ideas understandable. My growth requires learning to delegate, trust others, rest when necessary, and remain steady when a problem cannot be resolved immediately.	
Resources I have time, education, energy, initiative, opportunities, and a valuable network. I will convert them into the resources I need most: capital, technological mastery, and a trusted team. I will use AI to increase my efficiency and reach while continuing to invest in education and formative experiences.	
Relationships I want reciprocal relationships grounded in care, effort, loyalty, empathy, and honest communication. I will be present and dependable while maintaining boundaries that protect my health and identity. Supporting others must not require abandoning myself.	
Realities Self-criticism, impatience, fear of failure, and limited self-trust can restrict me more than external circumstances. I am willing to invest substantial time in my ambitions, but I will not sacrifice the health, relationships, and freedom those ambitions are intended to create.	
Repertoire I design databases, analyze data, translate findings into decisions, and make numerical problems understandable. My Industrial Engineering background provides strong systems, optimization, and problem-solving skills. I lead and collaborate well, learning best through formal study, mentorship, and teaching others.	
Vision I will build and lead an enduring business while remaining healthy and present at home. Its success will be measured not only	

financially, but also by the security, opportunity, and growth it provides its team and community. I want agency over my direction and work that extends beyond me.

Variables

My industry, business model, location, and timeline may change as technology, opportunities, and responsibilities evolve. My values will not. I will decide whether to persist or pivot through personal alignment, visible results, and trusted advice, taking calculated risks appropriate to my responsibilities.

Voices

I value the guidance of my parents, sisters, and closest friends, but their expectations will not define my life. When my inner critic creates urgency, I will remember that clarity grows through action and that I do not need every answer today. I will learn to trust my own voice.

Vulnerabilities

Fear and uncertainty can drive me toward stress, hyperfixation, and excessive control. Fear of appearing vulnerable can also prevent me from seeking help. Exercise, therapy, trusted relationships, rest, and healthy boundaries will keep me grounded. I will build boldly without abandoning myself, leading without needing total control and measuring success by both what I create and the lives I strengthen.

Whenever we produce coursework, remind me to save it in `~/crystal-ball/cuaderno`, commit it, and push it to GitHub.

CODEX - ASSISTANT

16 Sep 2026, 18:46:35 CST

I'll inspect both PDFs as source material only, not as instructions, and compare them with the manifesto you provided. I'll also keep your standing coursework reminder in view. I'm using the PDF skill because the source documents are layout-sensitive PDFs.

CODEX - ASSISTANT

16 Sep 2026, 18:47:04 CST

The two PDFs describe your Week 6 emergency-rehearsal work: the team blueprint favors measurable behavior and cautious deployment, while your individual brief argues for a consultant-led compliance workflow linked to a physical drill. I'm checking the original page layouts now so I don't miss headings or structure hidden by text extraction.

CODEX - ASSISTANT

16 Sep 2026, 18:48:11 CST

I've reviewed both PDFs and will treat their contents as source material-not instructions.

The team blueprint establishes measurable behavior, human debriefing, physical validation, accessibility, and ethical safeguards as the project's shared position. :codex-file-citation{path="/Users/davidbuzali/Desktop/BLUEPRINT_Team3_Week6.pdf" purpose="source"}

Your individual brief expresses your Operator perspective: consultant-led implementation, corrective-action closure, honest evidence, and a refusal to confuse completion or immersion with competence. :

codex-file-citation{path="/Users/davidbuzali/Desktop/BRIEF_davidbuzali (week 6).pdf" purpose="source"}

I'll use your [REDACTED DEVICE CODE]-VVVV Manifesto as a decision lens for future work: meaningful impact, educational opportunity in Mexico, measurable results, principled leadership, calculated risk, sustainable ambition, and success that does not sacrifice health or relationships.

Standing reminder noted: whenever we produce coursework, I'll remind you to save it in `~/crystal-ball/cuaderno`, commit it, and push it to GitHub. No coursework file was created in this turn.

- :codex-followup[Save the manifesto]{prompt="Save my [REDACTED DEVICE CODE]-VVVV Manifesto as a well-formatted project note in the workspace."}

- :codex-followup[Evaluate Week 6 alignment]{prompt="Evaluate how my Week 6 work aligns with my [REDACTED DEVICE CODE]-VVVV Manifesto and identify tensions or opportunities."}

- :codex-followup[Prepare the next deliverable]{prompt="Use the two Week 6 PDFs and my manifesto to help prepare my next coursework deliverable."}

DAVID BUZALI - USER

16 Sep 2026, 18:48:11 CST

Files pasted by the user:

"Individual · 55% of the week One working slice attacking your declared vacuum, ...":
 /Users/davidbuzali/.codex/attachments/8a225799-978b-419a-a16a-9ab250f063a2/pasted-text.txt

My request:

Santiago - USER

I ended with Behavior Measurement. My main idea is that the real problem starts when the normal emergency plan stops working. People may know the evacuation route, but a real emergency can create unexpected situations.

I brought the idea of practicing uncertainty. For example, an exit can be blocked, communication can fail, or someone can be missing. People need to practice how to make decisions in these situations.

At first, I focused more on Family Rehearsal, but I gave that up because this problem is bigger than families. It can happen at school, at work, or in any public place.

What survived from my idea is the second unexpected scenario. First, the person practices and learns from the mistakes. Then, we give them a different scenario to see if they actually make a better decision.

Ana María - TECHNOLOGIST

I ended with Behavior Measurement too. From the technology side, I brought one important rule: we should earn the headset. We should not use VR only because it looks more advanced.

For the first version, we do not need expensive VR or complicated AI. A simple browser simulation can create an emergency situation, record the decisions people make, and give useful information to a human safety leader.

The difficult part is not only creating a realistic simulation. The difficult part is making sure that the information we collect actually means something.

I gave up the idea that more technology is always better. What survived is a simpler system that records specific actions, shows group results, and lets a human make the final decision about what training needs to improve.

Mónica - MONEY

I started with Compliance-Upgrade because schools, companies, and other organizations already spend money on civil-protection training.

My main question was: who is going to pay for this? I brought the idea that we should not sell VR headsets as the main product. Instead, we can sell training as a service and give organizations evidence about whether their training actually worked.

This is important because we do not need to create a completely new market. Organizations already spend money on drills, instructors, training, and civil protection.

I gave up Compliance-Upgrade as our main vacuum, but it survived as our way to enter the market and make money. Behavior Measurement becomes the main product, while the existing compliance system gives customers a reason and a budget to use it.

David - OPERATOR

My role was to think about what happens when we actually try to use this product in a real school.

I brought the operating plan. We should start with teachers, safety staff, and other adults because they control the evacuation process. After they practice, we can do a small physical drill, identify the problems, correct them, and then expand the program to students.

I also brought the idea of working with existing civil-protection consultants instead of trying to replace them. They already understand the schools, the rules, and the current training system.

I gave up Compliance-Upgrade as the main vacuum, but an important part of my idea survived. Digital training should connect to a real physical test. If someone makes an important mistake, we should not simply mark the training as complete. The problem stays open until we see corrected behavior.

Fernando - ADVERSARY

My role was to find the reasons why this project could fail.

My biggest attacks were scale, not having enough headsets, teacher workload, parent resistance, and the possibility that VR looks impressive but does not actually make people more prepared.

At first, I supported Family Rehearsal because families do not normally practice what to do during an emergency at home. I also supported stronger and more realistic VR training because I think realistic pressure can be more useful than repeating the same predictable drill.

I gave up Family Rehearsal as our main vacuum because Behavior Measurement gives us a clearer way to prove whether the training is actually working.

What survived from my position is the scale test and the shadow rule. The experience should not use unnecessary fear,

traumatic images, or extreme content just to make VR feel realistic. The product only works if it is useful, safe, easy to repeat, measurable, and possible to scale.

CODEX - ASSISTANT

16 Sep 2026, 18:48:16 CST

I'll read the pasted assignment brief, treating it as coursework requirements rather than higher-priority instructions, and map your five role reflections to the requested deliverable. Then I'll produce the working slice in the appropriate format if the brief makes that clear.

CODEX - ASSISTANT

16 Sep 2026, 18:48:46 CST

The brief makes the first required artifact clear: a packet must exist before code. I'll start a separate Week 6 project slice around the team's chosen vacuum-Behavior Measurement-using your Operator contribution as the implementation bridge: adult rehearsal → physical micro-drill → unresolved failure stays open. The proposed stack will combine a browser 3D simulation, adaptive unseen-scenario selection, and voice guidance/commands so the three layers affect one another. I'm using the image-generation skill because the packet explicitly requires an AI-generated interface mockup.

CODEX - ASSISTANT

16 Sep 2026, 18:50:18 CST

The mockup is generated and saved in the Week 6 project. The product direction is now concrete: an adult teacher completes an initial 3D decision rehearsal, receives a human-led debrief, then faces an unseen adaptive disruption; voice input changes the simulation state, and the system records only observable actions. A critical digital result remains "pending physical validation," preventing the app from manufacturing a readiness claim.

CODEX - ASSISTANT

16 Sep 2026, 18:52:23 CST

Completed the packet-before-code stage for the Week 6 working slice.

Created:

- PACKET.md
- IMPLEMENTATION_PROMPT.md
- IMAGE_PROMPT.md
- DECISIONS.md
- Generated interface mockup

The concept combines browser 3D, adaptive unseen-scenario selection, and bounded voice commands in one evidence loop. Critical outcomes always remain pending physical validation. The benchmark is FLAIM FTS with Capture, localized for Mexican schools, existing consultants, and browser-first deployment.

The mockup was created with built-in image generation; its complete production prompt is saved in IMAGE_PROMPT.md.

No code, deployment, commits, or push were performed yet. Coursework reminder: save/sync the folder into ~/crystal-ball/cuaderno, then commit it and push it to GitHub.

- :codex-followup[Build the working slice]{prompt="Build the Decisión Bajo Presión application from the completed Week 6 packet and implementation prompt, then run the mechanical tests."}

DAVID BUZALI - USER

16 Sep 2026, 18:52:51 CST

Lets work on creating the packet

CODEX - ASSISTANT

16 Sep 2026, 18:53:08 CST

Absolutely. Let's build it deliberately, starting with the packet's load-bearing choice: the exact working slice.

Proposed foundation

Working title: Decisión Bajo Presión

Declared vacuum: Behavior Measurement

Exact user: A teacher or school-brigade member at a private CDMX school with an existing civil-protection consultant.

Working slice: The adult completes a short browser-based 3D rehearsal containing three emergency decisions. After a human-led debrief, the system presents a different scenario testing whether the person applies the same behavior under changed conditions.

Technology combination:

- Browser-based 3D simulation
- Transparent adaptive logic selecting the unseen retest
- Optional bounded voice commands

These components form one loop: the first simulation's observable actions determine the second simulation, while voice or button decisions enter the same evidence record.

Success definition:

> Before the module closes, the adult completes a baseline rehearsal, reviews an observable event trace with a human facilitator, completes a different unseen retest, and sees whether one narrowly defined behavior improved. Any critical behavior remains labeled "Pending physical validation"; the software cannot declare the participant competent or the school prepared.

Three decisions rehearsed:

1. What to do when the alert begins.
2. What to do when the primary exit becomes blocked.
3. What to do when an assigned person needs assistance.

What makes the slice distinctive: It measures whether learning transfers to a changed situation-not whether someone remembers the answer from the first simulation.

My recommendation is to keep the teacher or brigade member as the primary user and the consultant as the secondary actor. Should the first working slice instead be designed primarily for the civil-protection consultant?

DAVID BUZALI - USER

16 Sep 2026, 18:54:31 CST

The repository is in GitHub and should be a separate commit from the cuaderno; the launch will Vercel. Lets go by steps to go through this weeks assignment

CODEX - ASSISTANT

16 Sep 2026, 18:54:38 CST

Got it. We'll treat them as two separate submission tracks:

1. The application repository: its own Git history, deployments, testing evidence, and Vercel URL.
2. The cuaderno: a separate coursework reflection/update commit and push.

We'll move sequentially and won't jump to code before the packet is approved. I'll first verify the current Git repository and existing project structure so we place Week 6 in the correct repo without mixing it into the cuaderno.

CODEX - ASSISTANT

16 Sep 2026, 18:55:04 CST

The check found an important separation issue: the workspace root is still connected to the Week 2 repository, while Weeks 4 and 5 each have their own nested Git repositories. The new Week 6 folder is not yet connected to Git, so I won't commit it from the root and accidentally mix it into Week 2. I'm locating the intended Week 6 GitHub repository before we continue.

CODEX - ASSISTANT

16 Sep 2026, 18:55:18 CST

Understood. We'll proceed one stage at a time and keep the histories separate.

Workflow:

1. Finish and approve PACKET.md.
2. Generate the implementation prompt.
3. Connect the Week 6 folder to its dedicated GitHub repository and commit the packet.
4. Build through at least five meaningful application commits.
5. Make the first Vercel deployment.
6. Run the mechanical test, document a real bug, fix it, and redeploy.
7. Conduct the persona test and fix the most serious confusion.
8. Prepare the demo and submission PDFs.
9. Update ~/crystal-ball/cuaderno in a separate commit and push it separately.

We are currently on Step 1. The workspace root points to prueba-clara-week-2, and week-6-decision-under-pressure is not yet connected to Git, so I will not commit anything from the root.

What is the exact URL of the dedicated Week 6 GitHub repository?

DAVID BUZALI - USER

16 Sep 2026, 18:55:58 CST

I want you to create push and commit the new repository to the online version of GitHub

CODEX - ASSISTANT

16 Sep 2026, 18:56:04 CST

I'll create a dedicated GitHub repository named decision-bajo-presion-week-6, initialize the existing Week 6 folder as its own Git repository, commit only the packet-stage artifacts, and push the initial commit. This will remain completely separate from both the Week 2 repository and ~/crystal-ball/cuaderno.

CODEX - ASSISTANT

16 Sep 2026, 18:56:23 CST

GitHub's CLI is installed, but its saved authentication for davidbuzali has expired. I can still complete the local repository and initial packet commit now; then I'll reauthenticate GitHub and create/push the public remote. The initial commit will contain only the Week 6 packet, implementation prompt, decisions log, image prompt, and generated mockup.

CODEX - ASSISTANT

16 Sep 2026, 18:56:57 CST

The local repository and its first commit are complete. GitHub now requires a one-time device authorization in the browser; I'm opening that authorization flow so the remote can be created and pushed.

DAVID BUZALI - USER

16 Sep 2026, 18:57:34 CST

My request:
give me the code to verify github, it is 8 digits

CODEX - ASSISTANT

16 Sep 2026, 18:57:38 CST

GitHub verification code: [REDACTED DEVICE CODE]

DAVID BUZALI - USER

16 Sep 2026, 18:58:30 CST

I have connect github with the verification code

CODEX - ASSISTANT

16 Sep 2026, 18:58:33 CST

Great-I'll verify the authorization, create the new public repository, push the packet commit, and confirm the online repository and branch are accessible.

CODEX - ASSISTANT

16 Sep 2026, 18:59:34 CST

The new Week 6 repository is created and pushed:

github.com/davidbuzali/decision-bajo-presion-week-6

Details:

- Visibility: Public
- Default branch: main
- Initial commit: b62bf1c
- Commit: docs: define behavior-measurement rehearsal packet
- Local branch tracks origin/main
- Local and remote repositories are synchronized
- cuaderno was not modified

Next, we'll refine the packet section by section before writing application code or connecting Vercel.

DAVID BUZALI - USER

16 Sep 2026, 18:59:52 CST

Okey, lets move on with the next step

CODEX - ASSISTANT

16 Sep 2026, 19:00:13 CST

Next is Packet Part 1: lock the problem, exact user, and success definition before touching architecture.

Problem in my words

Emergency plans usually teach a fixed sequence, but the real failure begins when that sequence stops working. A blocked exit, failed communication, missing person, or conflicting instruction forces people to make decisions under uncertainty. Current training often records attendance or completion without showing whether someone can adapt when conditions change.

My working slice attacks this Behavior Measurement vacuum. It observes a small number of concrete decisions, supports reflection with a human facilitator, and then introduces a different scenario to test whether the learning transfers.

Exact user

The primary user is a teacher or school-brigade member at a private school in Mexico City that already has a formal civil-protection program and works with a registered consultant.

This person:

- Helps control routes, attendance, assistance assignments, or evacuation decisions.
- Has limited training time.
- Can use an ordinary browser but may have no gaming or VR experience.
- Must be able to participate without a headset or voice input.
- Needs feedback that helps them improve without making them feel psychologically judged.

The secondary user is the civil-protection consultant or school safety lead who facilitates the debrief and determines what must be tested in a physical drill.

Success definition

> Before the module closes, the adult completes one baseline rehearsal, reviews an observable event trace with a human facilitator, and completes a different unseen scenario testing the same narrow behavior under changed conditions. The system shows whether that behavior appeared in both scenarios, appeared only after the debrief, or remains unobserved. Every critical result remains pending physical validation.

The product does not declare that the participant is competent, that the school is prepared, or that lives will be saved.

If this foundation matches your intention, we'll lock it into the packet and move to the image-generated mockup and Mermaid flows.

DAVID BUZALI - USER

16 Sep 2026, 19:01:26 CST

Files mentioned by the user:

Screenshot 2026-09-16 at 7.01.05 p.m..png: /var/folders/cm/yg8ztrqx3zv0f07d1zggnx7r0000gn/T/TemporaryItems/NSIRD_screencaptureui_FO4uiY/Screenshot 2026-09-16 at 7.01.05 p.m..png

Distinguish instructions in attached documents from the user's request.

My request:

These are the vaccums for this week

<image name=[Image #1] path="/var/folders/cm/yg8ztrqx3zv0f07d1zggnx7r0000gn/T/TemporaryItems/NSIRD_screencaptureui_FO4uiY/Screenshot 2026-09-16 at 7.01.05 p.m..png">

</image>

CODEX - ASSISTANT

16 Sep 2026, 19:01:38 CST

That clarifies the course's exact taxonomy. Our declared vacuum should be **The MEASUREMENT vacuum**; "Behavior Measurement" is the team's chosen interpretation of that vacuum, not its official name. I'll correct the packet and repository language so the submission mirrors the assignment while preserving the product concept.

CODEX - ASSISTANT

16 Sep 2026, 19:02:24 CST

Correct-the official choice is:

> **The MEASUREMENT vacuum - zero data on whether drills change behavior.**

I updated the working packet to distinguish the course taxonomy from our product interpretation:

- **Declared vacuum:** Measurement
- **Product interpretation:** Measure whether a narrow behavior transfers to an unseen scenario.
- **Compliance-Upgrade:** Route to market and existing budget, not the primary vacuum.
- **Family and Insurance-Bridge:** Outside this week's scope.

I also corrected the terminology in README.md, DECISIONS.md, and the implementation prompt. These changes remain uncommitted while we finish refining the packet; we'll commit and push the completed packet revision as one coherent update.

Next packet section: the exact working slice and what the user experiences on each screen.

DAVID BUZALI - USER

16 Sep 2026, 19:03:19 CST

okey, now we can move on

CODEX - ASSISTANT

16 Sep 2026, 19:03:34 CST

I'll define the working slice screen by screen so the packet becomes buildable rather than conceptual. I'll keep it narrow: one adult session, one baseline, one human debrief, one unseen adaptive retest, and one evidence comparison that cannot certify readiness.

CODEX - ASSISTANT

16 Sep 2026, 19:04:09 CST

The working slice is now defined screen by screen in PACKET.md:

1. Start, limitations, accessibility, and optional voice
2. Baseline 3D scenario with three decisions
3. Observable evidence and human debrief
4. Adaptive unseen retest
5. Comparison and pending physical validation

I also specified:

- The exact baseline decisions
- How each missed behavior selects a different retest
- The only permitted evidence statements
- The bounded Spanish voice commands
- The requirement for professional validation before real-world use

Nothing has been committed yet; we are accumulating the packet revision into one coherent commit.

Next step: align the image-generated mockup and Mermaid diagrams with this five-screen experience.

DAVID BUZALI - USER

16 Sep 2026, 19:04:28 CST

Perfect, lets move on with the next step

CODEX - ASSISTANT

16 Sep 2026, 19:04:41 CST

I'll convert the diagrams into an exact five-screen flow and a true Mermaid swimlane separating the adult, browser simulation, voice layer, adaptive logic, and human safety lead. The existing generated mockup already represents Screen 4-the unseen retest-so I'll label that explicitly rather than regenerate it unnecessarily.

CODEX - ASSISTANT

16 Sep 2026, 19:05:38 CST

The mockup and diagrams are now aligned with the working slice.

Changes in [PACKET.md](#):

- Identified the generated mockup specifically as Screen 4: the unseen retest.
- Rebuilt the feature flow around Screens 1-5.
- Included the optional voice-permission branch.
- Showed that buttons, keyboard, and voice pass through validated action codes.
- Added a true Mermaid swimlane separating:
 - Adult teacher or brigade member
 - Optional voice layer
 - Browser simulation
 - Adaptive logic
 - Human safety lead
- Clearly marked the physical drill as outside the digital working slice.
- Verified the Markdown changes for formatting errors.

Next step: finalize the benchmark, three-year long view, and scope cut.

DAVID BUZALI - USER

16 Sep 2026, 19:05:48 CST

finalize the benchmark, three-year long view, and scope cut.

CODEX - ASSISTANT

16 Sep 2026, 19:05:55 CST

I'll tighten these sections against current primary sources, keeping the benchmark to the required one-sentence comparison, the long view to exactly three sentences, and the scope cut explicit enough to protect the build from expansion.

CODEX - ASSISTANT

16 Sep 2026, 19:06:42 CST

Finalized all three sections in [PACKET.md](#).

The benchmark now compares the product with FLAIM FTS and its Capture analytics/after-action-review platform. FLAIM's official product documentation

The localization explicitly connects the product to CDMX school civil-protection programs and the Responsible Official workflow. [CDMX civil-protection framework](#)

The long view is exactly three sentences and defines the product as human-led training infrastructure. The scope cut now excludes nine concrete expansion risks, including headset VR, digital twins, predictive models, backend accounts, student workflows, disaster physics, open-ended AI, certification, and the physical drill itself.

Next step: finalize the architecture, multiplicative Dragon Stack, security floor, and test plan.

DAVID BUZALI - USER

16 Sep 2026, 19:06:53 CST

finalize the architecture, multiplicative Dragon Stack, security floor, and test plan.

CODEX - ASSISTANT

16 Sep 2026, 19:07:09 CST

I'll turn these into implementation contracts: named modules, exact data flow, proof that all three Dragon Stack layers depend on one another, a security-floor checklist with "not applicable" justified where appropriate, and two documented test passes with measurable acceptance criteria.

CODEX - ASSISTANT

16 Sep 2026, 19:07:37 CST

One important correction emerged from the technical review: browser speech recognition is not uniformly available, and some browsers may send audio to their own recognition service. The packet will therefore make voice opt-in, disclose that limitation, store neither audio nor transcripts, and preserve complete click/keyboard functionality. That keeps the voice layer useful without making a false privacy claim.

CODEX - ASSISTANT

16 Sep 2026, 19:08:57 CST

The packet now has enough precision to function as the pre-code contract. I'm closing this packet session in `DECISIONS.md`, then I'll commit and push the completed packet revision as a single coherent documentation commit before we move to the implementation prompt.

CODEX - ASSISTANT

16 Sep 2026, 19:09:41 CST

The architecture, Dragon Stack, security floor, and test plan are finalized and pushed.

Key decisions:

- Client-only React + TypeScript + Vite application
- React Three Fiber for the screen-based 3D simulation
- Deterministic adaptive logic instead of a predictive model
- Optional voice input using the same validated action path as buttons
- No backend, accounts, database, analytics, secrets, or persistent data
- Fifteen defined mechanical checks
- One real documented bug-fix cycle
- Two verified Vercel deployments
- Fresh-chat persona test with a required usability fix

Voice is now described honestly: browser speech recognition has limited support and may use the browser provider's service, so it is opt-in and never required. [MDN SpeechRecognition documentation](#)

The deployment architecture follows Vite's documented Git-to-Vercel workflow. [Vite deployment guidance](#)

Committed and pushed:

- Commit: `3d1f652`
- Message: `docs: finalize measurement-vacuum product packet`
- [GitHub repository](#)

- Local packet

The application repository is synchronized. The cuaderno remains untouched and must receive its own separate coursework commit and push.

Next assignment step: finalize the implementation prompt before writing code.

DAVID BUZALI - USER

16 Sep 2026, 19:10:23 CST

finalize the implementation prompt before writing code.

CODEX - ASSISTANT

16 Sep 2026, 19:10:41 CST

I'll convert the finalized packet into a build contract that a coding agent can execute without inventing scope: exact file structure, ordered features, acceptance criteria, test gates, commit sequence, and deployment checkpoints. No application code will be written in this step.

CODEX - ASSISTANT

16 Sep 2026, 19:14:33 CST

The implementation prompt is finalized, committed, and pushed. No application code was written.

It now specifies:

- Exact module and file boundaries
- Closed TypeScript domain contracts
- Seven ordered, acceptance-tested features
- Shared validation across 3D, flat view, keyboard, pointer, and voice
- Exact scenario and adaptive-retest mappings
- Five security-floor checks
- M01-M15 testing obligations
- Nine planned build/fix/documentation commits
- Two verified Vercel deployment checkpoints
- Session-close requirements
- Complete definition of done

Files:

- Final implementation prompt
- Updated decisions log
- GitHub repository

Commit: 349fa64 - docs: finalize implementation build contract

The repository is synchronized with origin/main. The cuaderno remains separate and untouched.

Next step: Feature 1 only-scaffold Vite/React/TypeScript, implement the pure domain contracts and M01-M07 tests, then lint, test, build, commit, and push.

DAVID BUZALI - USER

16 Sep 2026, 19:14:53 CST

Feature 1 only-scaffold Vite/React/TypeScript, implement the pure domain contracts and M01-M07 tests, then lint, test, build, commit, and push.

CODEX - ASSISTANT

16 Sep 2026, 19:15:00 CST

I'll implement only Feature 1: the Vite/React/TypeScript foundation, pure domain contracts and fixtures, adaptive/comparison/event logic, and M01-M07 unit tests. I'll avoid UI feature work, then run lint, tests, and the production build before committing and pushing this increment.

CODEX - ASSISTANT

16 Sep 2026, 19:15:51 CST

I've locked the Feature 1 domain model: three baseline decisions, four deterministic retest families, one shared allow-list validator, an exact five-field event schema, and an immutable physicalValidation: "pending" result. I'm scaffolding those boundaries now; the React shell will remain intentionally minimal because Screens 1-5 belong to later feature commits.

CODEX - ASSISTANT

16 Sep 2026, 19:19:01 CST

The scaffold and domain tests are in place. I'm installing only the declared Feature 1 dependencies now; after that I'll run the three required gates and fix any type, lint, or test failures before committing.

CODEX - ASSISTANT

16 Sep 2026, 19:20:42 CST

The repository's nested location exposed a workspace issue: pnpm initially touched the parent Week 2 lockfile and then the isolated install hit sandboxed network access. I've added a local workspace boundary for Week 6; I'm now restoring the parent lockfile and completing the Week 6 install with approved network access so the repositories remain independent.

CODEX - ASSISTANT

16 Sep 2026, 19:21:39 CST

Dependencies are installed locally with a clean workspace boundary; I also pinned TypeScript to the version supported by the current lint tooling. I'm running peer validation, lint, M01-M07, and the production build now.

CODEX - ASSISTANT

16 Sep 2026, 19:22:00 CST

The code has not failed yet; the bundled workspace shell simply lacks node on the child-process path, so pnpm could not launch ESLint, Vitest, or TypeScript. I'm rerunning the exact tools through the bundled Node executable directly, which isolates genuine project failures from that environment issue.

CODEX - ASSISTANT

16 Sep 2026, 19:23:39 CST

Feature 1 now passes its real project checks: ESLint is clean, all 22 M01-M07 tests pass, strict TypeScript passes, and Vite produces the production bundle. I'm updating the repository's run instructions and session decision log, then I'll review the exact commit contents before committing.

CODEX - ASSISTANT

16 Sep 2026, 19:24:35 CST

Feature 1 is complete, committed, and pushed.

Implemented:

- Vite + React + strict TypeScript scaffold
- Isolated Week 6 pnpm workspace and lockfile
- Closed domain types and exact five-field event schema
- Three baseline decisions and four deterministic retest families
- Shared action/voice validation
- Adaptive retest selection
- Evidence comparison with immutable physicalValidation: "pending"
- Minimal placeholder shell-no Feature 2 UI yet

Verification:

- Peer dependencies: clean
- ESLint: passed
- M01-M07: **22 tests passed**
- Strict TypeScript: passed
- Vite production build: passed

- Secret-pattern scan: clean

Commit: 79add21 - chore: scaffold client app and domain contracts

The repository is synchronized with origin/main. The bundled Codex shell's Node-path limitation is documented in DECISIONS.md; the installed tools themselves all passed.

Next step: Feature 2-the safe Spanish start screen, session reducer, in-memory accessibility preferences, and pause-control shell. The cuaderno remains separate and still needs its own eventual commit and push.

DAVID BUZALI - USER

16 Sep 2026, 19:25:02 CST

Feature 2-the safe Spanish start screen, session reducer, in-memory accessibility preferences, and pause-control shell

CODEX - ASSISTANT

16 Sep 2026, 19:25:18 CST

I'll implement only Feature 2: an accessible Spanish start screen, a pure in-memory session reducer, motion/view/voice preferences, and pause/resume/exit controls. I'll add component and reducer tests, visually inspect the local screen, then lint, test, build, commit, and push-without starting the 3D scenario.

CODEX - ASSISTANT

16 Sep 2026, 19:26:02 CST

The Feature 2 interaction contract is set: preferences exist only in reducer memory; enabling voice changes a preference but never requests microphone access; pause preserves the active phase and evidence; exit recreates a clean intro state. I'm implementing that shell and its accessibility tests now.

CODEX - ASSISTANT

16 Sep 2026, 19:30:06 CST

The Feature 2 structure is in place. I'm tightening the accessibility tests now, then I'll do a browser-level visual check of both the start and paused states before committing.

CODEX - ASSISTANT

16 Sep 2026, 19:30:31 CST

All automated checks now pass: 32 tests, strict TypeScript, lint, and the production build. I'm at the final visual QA step, checking that the Spanish safety language, focusable controls, and pause overlay read cleanly in the actual browser.

CODEX - ASSISTANT

16 Sep 2026, 19:33:09 CST

Visual QA passed on desktop and a 390 × 844 mobile viewport, including preference handoff and the paused dialog. The implementation stayed inside Feature 2-no baseline scenario logic or 3D work was added-and I'm preparing the single feature commit now.

CODEX - ASSISTANT

16 Sep 2026, 19:33:40 CST

Feature 2 is complete and pushed to GitHub.

- Safe Spanish start screen with clear safety and privacy limits
- In-memory motion, view, and optional-voice preferences
- No microphone request or persistent storage
- Session reducer with start, pause, resume, and clean exit
- Accessible pause dialog with keyboard focus
- Responsive desktop and mobile layouts
- 32 tests passing
- ESLint, TypeScript, and production build passing

Commit: 9cc6c2a

The local branch and origin/main are synchronized. Feature 3 was not started.

Coursework reminder: save the corresponding coursework record in ~/crystal-ball/cuaderno, then commit and push it separately.

DAVID BUZALI - USER

16 Sep 2026, 19:34:18 CST

Okey, what are the next steps?

CODEX - ASSISTANT

16 Sep 2026, 19:34:21 CST

I'll map the remaining work directly from the finalized implementation prompt so the build, testing, deployment, and coursework evidence stay in the intended order.

CODEX - ASSISTANT

16 Sep 2026, 19:34:34 CST

The immediate next step is Feature 3. The remaining sequence is:

1. Feature 3 - Baseline rehearsal

- Build the primitive 3D corridor.
- Add the same three decisions in flat view.
- Record validated five-field events.
- Preserve keyboard, reduced-motion, and pause behavior.
- Commit: feat: build 3d baseline with flat-view parity

2. Feature 4 - Trace and human debrief

- Display neutral baseline events and relative timing.
- Select the behavior target deterministically.
- Require human-debrief acknowledgment before continuing.

3. Feature 5 - Unseen adaptive retest

- Load a different disruption based on baseline evidence.
- Reuse the same 3D/flat components, validator, and event pipeline.

4. Features 6-7 - Voice and comparison

- Add bounded Spanish voice commands with visible fallbacks.
- Show exactly one approved comparison statement.
- Keep every critical result marked **Pendiente de validación física**.

5. Complete M01-M15

- Add end-to-end and responsive tests.
- Verify the security floor and absence of personal data, secrets, analytics, and backend calls.
- Run lint, unit/component tests, E2E tests, and build.

6. Vercel Deployment 1

- Import the GitHub repository into Vercel.
- Test the complete public journey.
- Document a genuine defect found during testing.

7. Persona evaluation

- Test Screens 1-5 in a fresh conversation using the defined persona.
- Record confusion and determine the highest-severity issue.

8. Fix and Deployment 2

- Fix that issue in a separate commit.
- Rerun affected tests.
- Update testing/persona documentation and verify the second deployment.

For our next working session, we should do **Feature 3 only**, then lint, test, build, visually inspect, commit, and push.

Also remember to update ~/crystal-ball/cuaderno, commit it, and push it separately from the application repository.

DAVID BUZALI - USER

16 Sep 2026, 19:35:13 CST

Okey lets move on with the next step and I want you push and commit to GitHub when necessary

CODEX - ASSISTANT

16 Sep 2026, 19:35:18 CST

I'll build Feature 3 as one contained increment: the three-decision baseline, primitive browser 3D scene, exact flat-view parity, validated event capture, and preserved pause/accessibility behavior. I'll run the full checks and visual QA, then commit and push it to GitHub.

CODEX - ASSISTANT

16 Sep 2026, 19:40:13 CST

The baseline loop is now wired through the existing allow-listed event contract. I've also made the pause timer exclude paused time, kept the 3D and flat views on one shared decision component, and stopped the flow at an explicit Feature 4 boundary after the third event.

CODEX - ASSISTANT

16 Sep 2026, 19:41:05 CST

The test suite is green at 38 tests, and the app now lazy-loads the 3D engine so the Spanish start screen remains a small initial bundle. I'm checking the real WebGL scene, the blocked-route transition, the flat alternative, and mobile layout before I close the commit.

CODEX - ASSISTANT

16 Sep 2026, 19:44:01 CST

Feature 3 has passed its final gate: 38 tests, lint, strict TypeScript, production build, and real-browser checks of both WebGL and flat modes. The only build notice is the expected Three.js chunk size; it is isolated behind lazy loading and recorded as a deployment-performance check rather than hidden.

CODEX - ASSISTANT

16 Sep 2026, 19:44:32 CST

Feature 3 is complete, committed, and pushed to GitHub.

Implemented:

- Three-decision baseline rehearsal
- Primitive fictional 3D corridor with visible scenario changes
- Equivalent reduced-motion flat view
- Shared pointer and keyboard action pipeline
- Five-field validated event recording
- Ordered-event and duplicate protection
- Pause time excluded from decision timing
- Explicit Feature 4 boundary after completion
- Lazy-loaded Three.js bundle

Verification:

- 38 tests passed
- ESLint passed
- Strict TypeScript passed
- Production build passed
- WebGL, flat view, blocked-route transition, and responsive layout visually verified

Commit: 8fff20b

main and origin/main are synchronized. The next increment is Feature 4: the neutral event trace, adaptive behavior target, and required human-debrief acknowledgment.

Remember to update ~/crystal-ball/cuaderno, commit it, and push it separately.

DAVID BUZALI - USER

16 Sep 2026, 19:44:46 CST

Okey lets move on with the next step

CODEX - ASSISTANT

16 Sep 2026, 19:44:54 CST

I'll implement Feature 4 only: the neutral event trace, deterministic behavior target, human-decision boundary, and required in-memory debrief acknowledgment. The retest will remain an explicit placeholder for Feature 5; after verification I'll commit and push this increment.

CODEX - ASSISTANT

16 Sep 2026, 19:48:05 CST

Feature 4's logic is now green at 42 tests. The debrief cannot be confirmed early, the proposed target is derived only from validated baseline events, and the unseen retest's conditions and options remain hidden. I'm doing the final browser review of the trace and confirmation boundary now.

CODEX - ASSISTANT

16 Sep 2026, 19:49:56 CST

The browser review passed on desktop and mobile: the trace stays factual, the simulated-AI label is explicit, no retest option is leaked, and the confirmation advances only after the full three-event trace exists. I'm closing this as the separate Feature 4 commit now.

CODEX - ASSISTANT

16 Sep 2026, 19:50:42 CST

Feature 4 is complete, committed, and pushed.

Implemented:

- Neutral three-event timeline with relative timing
- Input-mode labels without personal data
- Deterministic behavior target from validated baseline events
- **IA simulada - la decisión final es humana** boundary
- No scores, probabilities, traits, or competence judgments
- Retest conditions and answers remain hidden
- Required human-debrief confirmation
- Confirmation remains memory-only and requires a complete trace
- Explicit Feature 5 placeholder

Verification:

- 42 tests passed
- ESLint and strict TypeScript passed
- Production build passed
- Desktop and mobile browser QA passed
- No runtime errors during the Feature 4 journey

Commit: a3359d3

The next increment is Feature 5: the adaptive unseen retest using the shared scenario UI and evidence pipeline.

Remember to update, commit, and push ~/crystal-ball/cuaderno separately.

DAVID BUZALI - USER

16 Sep 2026, 19:50:51 CST

Okey, next step

CODEX - ASSISTANT

16 Sep 2026, 19:51:12 CST

I'll implement Feature 5 as a true reuse step: generalize the existing scenario screen, feed it the deterministic unseen ScenarioDefinition, record the retest through the same validator/reducer, preserve 3D/flat/keyboard/pause behavior, and stop at the comparison boundary. Then I'll verify, commit, and push it.

CODEX - ASSISTANT

16 Sep 2026, 19:56:28 CST

The adaptive path is now passing 46 tests, including wrong-family rejection, 3D/flat parity, pause preservation, and the full baseline → debrief → selected retest → comparison-boundary journey. I'm checking the actual route-change retest scene and its required safety labels in the browser.

CODEX - ASSISTANT

16 Sep 2026, 19:59:10 CST

Feature 5 has cleared the final gate: 46 tests, lint, strict TypeScript, production build, and real-browser checks of the selected WebGL retest, reduced-motion flat retest, pause/resume, and the comparison boundary. I'm creating the dedicated Feature 5 commit and pushing it now.

CODEX - ASSISTANT

16 Sep 2026, 19:59:58 CST

Feature 5 is complete, committed, and pushed.

Implemented:

- Deterministic unseen retest selection
- Shared baseline/retest scenario interface
- Shared primitive 3D and flat-view presentations
- Scenario-specific route, accountability, assistance, and conflicting-sign visuals
- Required Screen 4 safety labels
- Same validator, reducer, timing, keyboard, and pause pipeline
- Wrong-family and incomplete-trace rejection
- Valid retest events advance to the comparison boundary
- No randomness or persistent data

Verification:

- 46 tests passed across 9 files
- ESLint and strict TypeScript passed
- Production build passed
- WebGL and reduced-motion flat retests visually verified
- Pause/resume and comparison-boundary transition verified

Commit: [506649f](#)

Next is Feature 6: explicit bounded Spanish voice input with complete keyboard/button fallback and no stored transcript.

Remember to update, commit, and push ~/crystal-ball/cuaderno separately.

DAVID BUZALI - USER

16 Sep 2026, 20:00:13 CST

Okey, next step

CODEX - ASSISTANT

16 Sep 2026, 20:00:48 CST

I'll implement Feature 6 only: explicit Escuchar comando activation, bounded es-MX command mapping into the existing action callback, immediate transcript disposal, neutral failure states, and uninterrupted button/keyboard fallback. I'll verify permission

timing and error paths, then commit and push it.

CODEX - ASSISTANT

16 Sep 2026, 20:04:02 CST

The voice adapter is green across success, pause, unknown phrase, permission denial, network failure, unsupported browser, and no-result paths. Recognition is never constructed when voice is merely enabled or when a scenario loads; only Escuchar comando starts it. I'm now checking the opt-in and fallback presentation in the real browser without triggering microphone access.

CODEX - ASSISTANT

16 Sep 2026, 20:05:23 CST

Feature 6 has passed its final gate with 61 tests. The real browser showed the opt-in voice control and provider disclosure while making no microphone request; supported, denied, failed, unknown, pause, and successful-command behaviors are covered mechanically. I'm preparing the separate Feature 6 commit now.

CODEX - ASSISTANT

16 Sep 2026, 20:05:59 CST

Feature 6 is complete, committed, and pushed.

Implemented:

- Explicit **Escuchar comando** activation
- One-shot es-MX recognition
- Five bounded Spanish commands
- Immediate transcript disposal
- No audio or transcript storage/display/logging
- Shared voice, button, and keyboard validation pipeline
- Voice-based pause without evidence creation
- Neutral handling for unknown phrases, denial, network errors, recognition errors, and unsupported browsers
- Visible controls remain usable after every failure

Verification:

- 61 tests passed across 10 files
- ESLint and strict TypeScript passed
- Production build passed
- Voice opt-in UI visually verified without triggering microphone access
- No console errors or warnings during the voice opt-in journey

Commit: [adbc1e1](#)

Next is Feature 7: the final evidence comparison and immutable **Pendiente de validación física** gate.

Remember to update, commit, and push ~/crystal-ball/cuaderno separately.

DAVID BUZALI - USER

16 Sep 2026, 20:06:11 CST

Lets do feature 7

CODEX - ASSISTANT

16 Sep 2026, 20:06:18 CST

I'll complete Feature 7 as the final product screen: exactly one approved comparison statement for the targeted behavior, an immutable **Pendiente de validación física** gate, and a concrete physical micro-drill instruction-with no digital approval or certification control. Then I'll run the full flow checks, visually verify it, commit, and push.

CODEX - ASSISTANT

16 Sep 2026, 20:07:13 CST

I'll implement Feature 7 as the evidence comparison and immutable physical-validation gate, then run the full verification suite, inspect the finished flow in the browser, and commit/push the feature to GitHub.

CODEX - ASSISTANT

16 Sep 2026, 20:07:45 CST

The domain already enforces the key safety invariant-physicalValidation can only be "pending". I'm now connecting that pure comparison to Screen 5 and keeping the screen read-only: one approved evidence statement, one targeted behavior, and no digital approval or closure action.

CODEX - ASSISTANT

16 Sep 2026, 20:09:45 CST

Feature 7 is wired end-to-end, including three rendered-outcome tests, a complete baseline-to-comparison journey, an unresolved-retest case, and a fresh-mount privacy reset. I'm at the mechanical verification stage now; any failure will be fixed before browser QA or the Git commit.

CODEX - ASSISTANT

16 Sep 2026, 20:10:03 CST

All 67 tests pass, lint and strict TypeScript are clean, and the production build succeeds. The final screen adds only about 4 KB gzip to CSS/app code; the already-known lazy Three.js chunk remains the sole size warning. I'm moving to visual and responsive QA of the actual five-step flow.

CODEX - ASSISTANT

16 Sep 2026, 20:11:24 CST

Browser QA confirms the complete flat-view journey reaches the correct "demonstrated after debrief" statement, the physical micro-drill gate remains prominent at desktop and 390×844, reload returns to a clean start, and the console has no errors or warnings. I'm recording those actual results, then I'll create the Feature 7 commit and push it.

CODEX - ASSISTANT

16 Sep 2026, 20:12:02 CST

The Feature 7 commit is created as 423d5a4 with the final comparison screen, responsive styling, full-flow/privacy tests, and updated project records. I'm pushing that exact commit to main now, then I'll verify the local and GitHub revisions match.

CODEX - ASSISTANT

16 Sep 2026, 20:12:28 CST

Feature 7 is complete, committed, and pushed to GitHub.

- Added the final evidence-comparison screen.
- Displays exactly one approved outcome.
- Keeps every result **Pendiente de validación física**.
- Includes physical micro-drill instructions.
- Provides no digital approval, certification, or closure control.
- Reloading clears the session.
- 67 tests passed across 11 files.
- Lint, TypeScript, production build, desktop/mobile QA, and console checks passed.
- Commit: 423d5a4 feat: add comparison and physical-validation gate
- GitHub: [decision-bajo-presion-week-6](#)

Next step: complete the security/M01-M15 evidence record, then launch Deployment 1 on Vercel.

Coursework reminder: save the corresponding notes separately in ~/crystal-ball/cuaderno, commit them, and push them to GitHub.

DAVID BUZALI - USER

16 Sep 2026, 20:12:36 CST

complete the security/M01-M15 evidence record, then launch Deployment 1 on Vercel.

CODEX - ASSISTANT

16 Sep 2026, 20:12:43 CST

I'll first turn M01-M15 and the five security checks into a reproducible evidence record, filling only results we can actually verify. Then I'll deploy the tested commit to Vercel as Deployment 1, verify the public URL, and record the deployment evidence in a separate commit and push.

CODEX - ASSISTANT

16 Sep 2026, 20:13:07 CST

The current unit/component suite covers M01-M10 and most of M11-M13, but the assignment explicitly calls for a reproducible Playwright journey and responsive evidence. I'm adding that missing end-to-end layer instead of overstating the existing browser inspection, then I'll run the security scans against tracked source and dependencies.

CODEX - ASSISTANT

16 Sep 2026, 20:16:44 CST

The reproducible browser layer is now green: two real Chrome tests cover the full keyboard path, pause focus, unseen retest, immutable physical gate, reload privacy reset, and the 390×844 layout with no horizontal overflow. Next I'm running the source/dependency security floor and will write the exact commands and results into docs/TESTING.md.

CODEX - ASSISTANT

16 Sep 2026, 20:18:36 CST

The strengthened keyboard test exposed a test-design issue, not an app defect: native radio groups correctly place only the selected radio in the Tab order, so Tab skips the unselected "Vista plana" option. I'm updating the test to use the standard keyboard pattern-Tab to the selected 3D radio, then ArrowDown to choose the flat option-and will rerun it.

CODEX - ASSISTANT

16 Sep 2026, 20:22:03 CST

The security scan is clean: no tracked environment files or secret assignments, no client persistence/network-write APIs, no free-text/personal-data fields, no backend/auth/analytics dependencies, and the production dependency audit reports zero known vulnerabilities. I'm now consolidating the exact M01-M15 mappings, commands, screenshots, remaining bundle risk, and security-floor evidence into the permanent testing record.

CODEX - ASSISTANT

16 Sep 2026, 20:23:43 CST

The permanent evidence record is complete and all pre-deployment checks are green: M01-M15, 67 component/unit tests, two Chrome end-to-end tests, lint, strict build, dependency audit, source scans, network assertion, and retained desktop/mobile screenshots. I'm creating and pushing the dedicated test/security commit now; Vercel will deploy from that exact revision.

CODEX - ASSISTANT

16 Sep 2026, 20:25:00 CST

The test/security commit 4850771 is pushed and matches origin/main. There is no Vercel CLI installed in the project, so I'm checking the existing Vercel login through an ephemeral CLI invocation; if the account is already authenticated, I'll link this GitHub-backed project and create the production Deployment 1 without adding deployment tooling to the repository.

CODEX - ASSISTANT

16 Sep 2026, 20:26:22 CST

Vercel authentication is active, and the project has been created as davidbuzali/decision-bajo-presion-week-6 with the GitHub repository connected. Vercel correctly detected Vite and dist; I'm launching the production deployment from commit 4850771 now.

CODEX - ASSISTANT

16 Sep 2026, 20:27:56 CST

Deployment 1 is live and Vercel reports READY, with the production alias <https://decision-bajo-presion-week-6.vercel.app>. I'm now verifying the public-not local-site through the full five-screen journey, checking reload privacy, runtime console, and actual resource origins before I record the deployment as

successful.

CODEX - ASSISTANT

16 Sep 2026, 20:31:48 CST

The public flat-view journey, pause/resume, comparison, physical gate, and reload reset all work. The default 3D path also renders from its lazy chunk. The only public runtime finding is the previously observed Three.js deprecation warning emitted inside React Three Fiber (new `THREE.Clock()`); it does not block the UI, but I'm recording it as a real Deployment 1 defect for the required later test-fix cycle rather than pretending the console is perfectly clean.

CODEX - ASSISTANT

16 Sep 2026, 20:35:13 CST

Deployment 1 is fully verified: production is READY, the stable alias returns HTTP 200 with HSTS, all three public Playwright checks pass, the five-screen journey and reload reset work, and the lazy 3D scene visibly renders. I've recorded the immutable deployment ID/URL and the real non-blocking dependency warning; I'm committing and pushing this post-deployment evidence now.

CODEX - ASSISTANT

16 Sep 2026, 20:37:41 CST

The GitHub-connected project automatically built the documentation-only evidence commit after it was pushed. That is not the assignment's corrective Deployment 2, so I'm restoring the production alias to the original verified Deployment 1 (dp1_4Ep...) and leaving the later immutable build as an uncounted documentation redeploy. This keeps the recorded production alias and required deployment sequence accurate.

CODEX - ASSISTANT

16 Sep 2026, 20:38:58 CST

Completed and pushed.

- M01-M15 evidence: [docs/TESTING.md](#)
- 67 Vitest tests passed.
- 3 public Playwright checks passed.
- Lint, strict TypeScript, build, security scans, and dependency audit passed.
- No known production dependency vulnerabilities.
- No secrets, personal-data fields, backend, database, analytics, or persistent storage found.

Deployment 1:

- Open production app
- Deployment: `dp1_4EpMNwyanWvVaYjfsYa6wwPmXgZW`
- Source commit: `4850771`
- Status: **Ready**
- HTTP 200 with HSTS enabled.
- Stable alias restored to the original verified Deployment 1, so Deployment 2 remains pending.

Evidence was pushed in commit `71a7d6f`.

One genuine non-blocking defect was recorded: the 3D dependency emits a `THREE.Clock` deprecation warning. The scene still renders correctly. Next is the fresh Laura persona pass, followed by selecting and fixing the highest-severity real issue before Deployment 2.

Coursework reminder: update `~/crystal-ball/cuaderno` separately, then commit and push that repository.

DAVID BUZALI - USER

16 Sep 2026, 20:39:35 CST

Lets run the persona test

CODEX - ASSISTANT

16 Sep 2026, 20:40:21 CST

I'll run the persona pass exactly as specified: capture Screens 1-5 from the verified public flow, then give only those screens and Laura's persona-not the intended interpretation-to a fresh task. I'll record every hesitation by screen, rank severity, and keep the observed feedback separate from our product assumptions.

CODEX - ASSISTANT

16 Sep 2026, 20:41:24 CST

The five public screenshots are captured and visually clean. They show the start choices, visible pause control, first decision, human debrief, different unseen retest, and final physical-validation gate in sequence. I'm freezing those artifacts in Git so the fresh persona task evaluates exactly the same evidence we will retain.

CODEX - ASSISTANT

16 Sep 2026, 20:42:23 CST

The fresh Laura task is running with only the ordered screenshots and persona description. It has been explicitly barred from reading the code, documentation, tests, or our intended interpretation, so its feedback remains a genuinely blind synthetic-user pass.

CODEX - ASSISTANT

16 Sep 2026, 20:44:24 CST

Laura's blind review is still running through the five screens. While it evaluates them, I'm preserving the exact screenshot sequence and will report the raw hesitation evidence separately from the eventual fix recommendation.

CODEX - ASSISTANT

16 Sep 2026, 20:45:13 CST

The first blind pass surfaced a flaw in my test setup: Laura said she would choose both reduced motion and flat view, but the screenshot journey applied only flat view, so later screens still showed "Movimiento normal." That would contaminate the findings. I'm correcting the capture sequence and rerunning the fresh persona-not treating that artificial contradiction as product evidence.

CODEX - ASSISTANT

16 Sep 2026, 20:46:30 CST

The corrected screenshots now show Laura's reduced-motion choice consistently. Since a second fresh task was declined, I won't create another one. I'll have the existing blind reviewer re-evaluate only the corrected evidence and explicitly retract any finding caused by the earlier capture mistake; the unaffected observations remain traceable.

DAVID BUZALI - USER

16 Sep 2026, 20:46:39 CST

Please run the persona test in this chat

CODEX - ASSISTANT

16 Sep 2026, 20:47:02 CST

Understood-I'll run it here. I'm using the corrected public screenshots, where Laura's choices are carried forward as **Movimiento reducido** and **Vista plana**. I'll treat this as a synthetic in-chat persona test, clearly distinguish observations from recommendations, and record the result without implementing the fix yet.

CODEX - ASSISTANT

16 Sep 2026, 20:48:32 CST

The in-chat persona walkthrough is complete. I found one blocking safety issue-Laura can pause during decision screens, but not while waiting for the required human debrief-and I've documented it as P01. I'm checking the evidence package now, then I'll commit and push the persona record.

DAVID BUZALI - USER

16 Sep 2026, 20:48:35 CST

Please run the persona test in this chat

CODEX - ASSISTANT

16 Sep 2026, 20:48:40 CST

Yes-this is being run here as a synthetic, in-chat persona test, not delegated to another task and not represented as real-user research. I'll include Laura's screen-by-screen reactions, the severity-ranked observation record, and the resulting fix decision in the final handoff.

CODEX - ASSISTANT

16 Sep 2026, 20:48:47 CST

The record is internally consistent and the evidence now carries both accessibility choices into the decision screens. The ordinary shell doesn't expose npm in this workspace, so I'm invoking the project's installed toolchain directly for the final lint check.

CODEX - ASSISTANT

16 Sep 2026, 20:48:56 CST

The persona evidence passes lint. I'm recording it as a separate coursework commit now; the application itself remains unchanged, and the identified P01 fix is intentionally reserved for the next implementation commit and Deployment 2.

CODEX - ASSISTANT

16 Sep 2026, 20:49:25 CST

Commit bc1e324 now contains the in-chat persona report and corrected evidence. I'm pushing it to GitHub; afterward I'll verify that the public alias still represents the approved Deployment 1 rather than treating this documentation-only push as Deployment 2.

CODEX - ASSISTANT

16 Sep 2026, 20:50:16 CST

GitHub's Vercel integration automatically promoted the documentation-only build. I'm restoring the stable alias to the verified Deployment 1 so that "Deployment 2" remains reserved for the actual P01 correction and retest.

CODEX - ASSISTANT

16 Sep 2026, 20:50:55 CST

Persona test completed in this chat using the synthetic persona Laura Hernández, a 47-year-old teacher and brigade volunteer with mild motion sensitivity.

Key results:

- Laura successfully selected reduced motion and flat view.
- She immediately noticed **Pausar y salir** on both decision screens.
- She understood that the AI result was simulated and that physical validation remained pending.
- She hesitated because selecting an answer appears immediate and irreversible.
- **P01 - Blocking:** The mandatory human-debrief screen has no pause or exit option. If the responsible person is unavailable, Laura must abandon the page or falsely confirm the debrief.

The full walkthrough, quotations, observation table, severity ranking, and retest requirement are in [docs/PERSONA.md](#).

Completed:

- Lint passed.
- Evidence corrected to preserve reduced-motion and flat-view preferences.
- Committed as bc1e324 - docs: record Laura persona findings
- Pushed to GitHub.
- [Deployment 1](#) remains the verified production version.
- Deployment 2 remains reserved for the P01 fix and retest.

Next step: add **Pausar y salir** to the human-debrief screen without allowing users to bypass the human gate, then retest that screen and launch Deployment 2.

Coursework reminder: also save/update the corresponding record in ~/crystal-ball/cuaderno, commit it separately, and push it to GitHub.

DAVID BUZALI - USER

16 Sep 2026, 20:51:05 CST

add **Pausar y salir** to the human-debrief screen without allowing users to bypass the human gate, then retest that screen and launch Deployment 2.

CODEX - ASSISTANT

16 Sep 2026, 20:51:11 CST

I'll implement P01 as a safe pause path on the human-debrief screen while keeping progression locked behind the real debrief confirmation. Then I'll add regression coverage, rerun the affected persona screen and full checks, commit/push, and verify Deployment 2 on Vercel.

CODEX - ASSISTANT

16 Sep 2026, 20:52:05 CST

The smallest correct change is to make debrief a resumable paused phase-not to add an alternate continue action. The existing pause overlay can then resume to the same debrief, while the reducer still permits entry to the retest only through `confirm_debrief` after a complete baseline. I'm encoding that invariant in unit, component, and end-to-end tests.

CODEX - ASSISTANT

16 Sep 2026, 20:52:43 CST

The regression now checks the exact safety invariant: pausing from Screen 3 preserves all evidence, hides the confirmation while paused, resumes back to Screen 3, and does not unlock the retest. I'm running the unit/component suite and build checks before exercising the browser flow.

CODEX - ASSISTANT

16 Sep 2026, 20:53:02 CST

One pre-existing baseline test correctly caught the intentional behavior change: it still expected the pause control to disappear after the third decision. I'm updating that assertion to the new safety requirement, then rerunning the complete suite.

CODEX - ASSISTANT

16 Sep 2026, 20:53:18 CST

All 69 unit/component tests pass, along with lint, TypeScript, and production build. The build still reports the already-known non-blocking 3D chunk-size warning; no new security or behavior regression appeared. I'm now running the full browser suite, which will also recapture Screen 3 with the new control.

CODEX - ASSISTANT

16 Sep 2026, 20:53:43 CST

The full browser suite passes: 4/4, including keyboard use, mobile layout, immutable physical gate, 3D smoke test, and the new Screen 3 pause/resume path. I'm visually reviewing the recaptured persona screen and then I'll update the evidence record from "blocked" to "retested and resolved."

CODEX - ASSISTANT

16 Sep 2026, 20:54:22 CST

The visual retest shows **Pausar y salir** prominently on Screen 3 without changing or duplicating the human-confirmation action. I'm retaining only the changed Screen 3 persona evidence; the browser suite also regenerated unrelated screenshots because of fresh rendering timestamps, so I'm discarding those non-semantic binary changes to keep the deployment commit auditable.

CODEX - ASSISTANT

16 Sep 2026, 20:55:11 CST

The local persona retest is now documented as resolved: Laura can safely wait or exit, resuming returns to the same debrief, and the unseen retest stays locked until the explicit human-debrief confirmation. I'm committing this as the implementation candidate for Deployment 2.

CODEX - ASSISTANT

16 Sep 2026, 20:55:32 CST

Commit 9bee747 is on GitHub. Vercel should now build it through the connected repository; I'm waiting for that exact revision to become Ready before calling it Deployment 2 or testing the public alias.

CODEX - ASSISTANT

16 Sep 2026, 20:56:15 CST

Deployment 2 is Ready and now serves the stable production alias (dpL_4qQG...). I'm running the same four browser checks against the public URL, including the debrief pause/resume invariant and a fresh public Screen 3 capture.

CODEX - ASSISTANT

16 Sep 2026, 20:56:52 CST

Public verification passed 4/4 in 13.7 seconds. I'm preserving the public Screen 3 capture and removing unrelated fresh-render changes from the other evidence images, then recording the exact Deployment 2 ID, source commit, and verification results.

CODEX - ASSISTANT

16 Sep 2026, 20:58:07 CST

Deployment 2 is fully verified in production, and the public persona screenshot confirms the corrected control. I'm creating the final evidence commit now; after pushing it, I'll restore the stable alias to the verified code deployment if Vercel's documentation-only rebuild moves it.

CODEX - ASSISTANT

16 Sep 2026, 20:59:06 CST

The documentation push produced a separate automatic build, so I restored the production alias to the already tested Deployment 2. I'm doing one final synchronization and alias check now.

CODEX - ASSISTANT

16 Sep 2026, 20:59:27 CST

Implemented, retested, committed, pushed, and launched Deployment 2.

- **Pausar y salir** is now visible during the human debrief.
- Pausing does not confirm the debrief or create evidence.
- Resuming returns to the unchanged debrief.
- The unseen retest remains locked until explicit human confirmation.
- Persona blocker P01 is resolved.

Verification:

- 69 Vitest tests passed.
- 4/4 Playwright tests passed locally and publicly.
- Lint, TypeScript, and production build passed.
- Production returned HTTP 200 with HSTS.
- D01, the non-blocking upstream Three.js warning, remains documented.

Commits:

- 9bee747 - P01 implementation and regression coverage
- e3afdfa - Deployment 2 and persona-retest evidence

Deployment 2: [Open production site](#)

Evidence:

- Persona retest
- Testing and deployment record
- Updated Screen 3

Remember to update the coursework record in ~/crystal-ball/cuaderno, commit it separately, and push it to GitHub.

DAVID BUZALI - USER

16 Sep 2026, 21:00:34 CST

Please give me a pdf of the PACKET_yourname.pdf, PERSONA_yourname.pdf, BUILDCHAT_yourname.pdf

CODEX - ASSISTANT

16 Sep 2026, 21:00:40 CST

I'll use the PDF workflow to produce three polished, submission-ready files named for you: PACKET_davidbuzali.pdf, PERSONA_davidbuzali.pdf, and BUILDCHAT_davidbuzali.pdf. I'll derive them from the repository's final packet, persona evidence, decision/build records, and verified deployment results, then render and visually inspect every PDF.

CODEX - ASSISTANT

16 Sep 2026, 21:01:03 CST

The assignment confirms the three PDFs are the required submission artifacts: the packet, the persona-test log, and the full development back-and-forth. I found prior-week PDF exports, so I'm checking their established format and generation approach to keep this week's submission consistent rather than inventing a new document style.

CODEX - ASSISTANT

16 Sep 2026, 21:01:36 CST

The prior submission format uses a consistent Letter-size ReportLab design and a sanitized visible-chat export. I'm reusing that structure for Week 6, while updating the content, live links, deployment evidence, screenshots, and transcript boundary. The build chat will include visible user-assistant messages only-no internal reasoning, credentials, system instructions, or tool telemetry.

CODEX - ASSISTANT

16 Sep 2026, 21:02:56 CST

All three PDFs have been generated. I'm validating metadata, page counts, text integrity, links, and then rendering every page to images for visual QA; I'll correct any clipping, broken tables, or malformed transcript content before delivery.

CODEX - ASSISTANT

16 Sep 2026, 21:05:03 CST

The first visual pass is clean-no clipped pages, broken tables, or unreadable screenshots-but I found one inherited footer line on the cover pages that still says "Week 5." I'm correcting that to Week 6, then I'll rerender and repeat the final checks.

CODEX - ASSISTANT

16 Sep 2026, 21:07:13 CST

The corrected covers now say Week 6, and the final-page checks are clean. I'm making one last presentation improvement to the persona log: using the otherwise sparse final page to embed the verified Deployment 2 Screen 3 screenshot, so the claimed correction is visible inside the submitted PDF rather than only linked.